

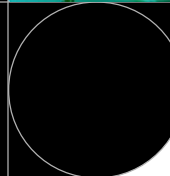
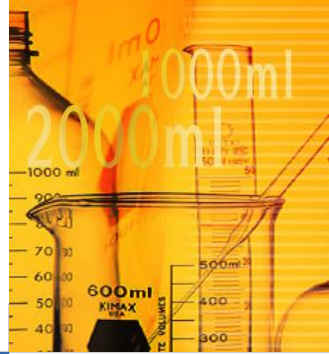
Machine learning

Chapter 10

Clustering

Sejong Oh

Bio Information Technology Lab.



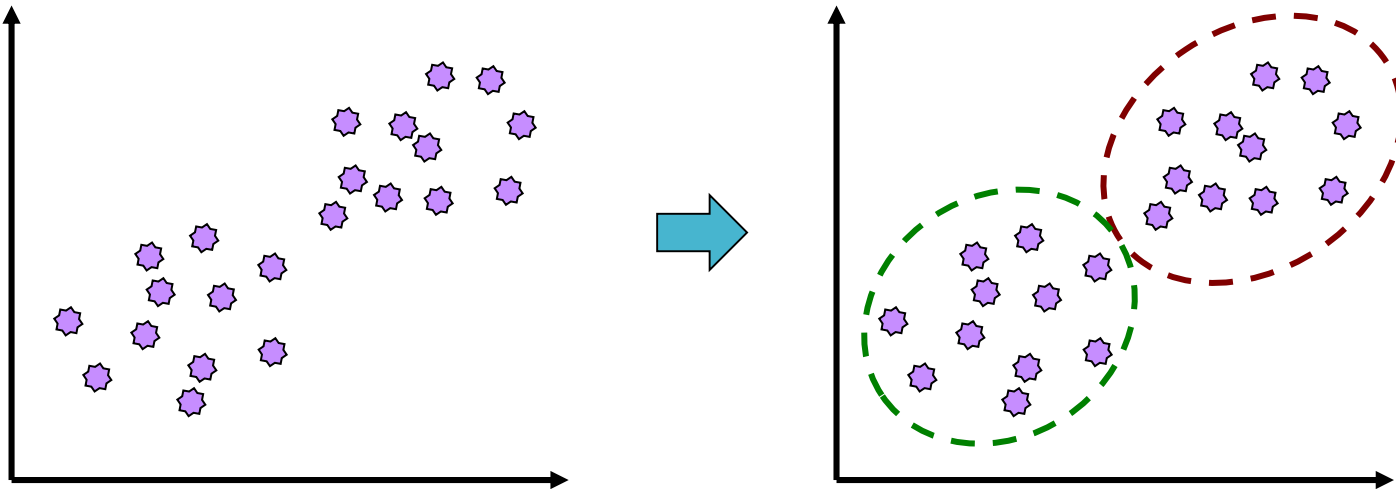
Contents



- Summary
- K-means clustering
- Hierarchical clustering

- Clustering

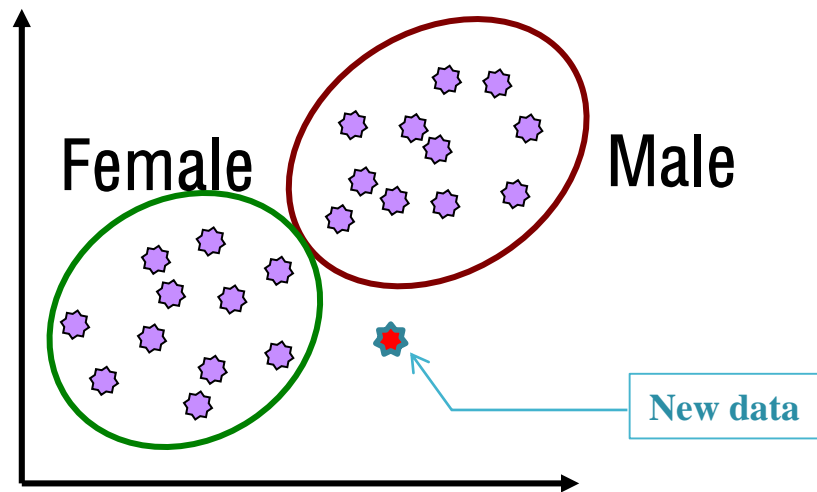
- Grouping target data into some category
- Data in same group has similar characteristics
- Group points into clusters based on how “near” they are to one another
- Unsupervised learning



Summary

- Classification

- Classify new data into one of known category.
- The category has “label”
- Supervised learning



- Clustering example

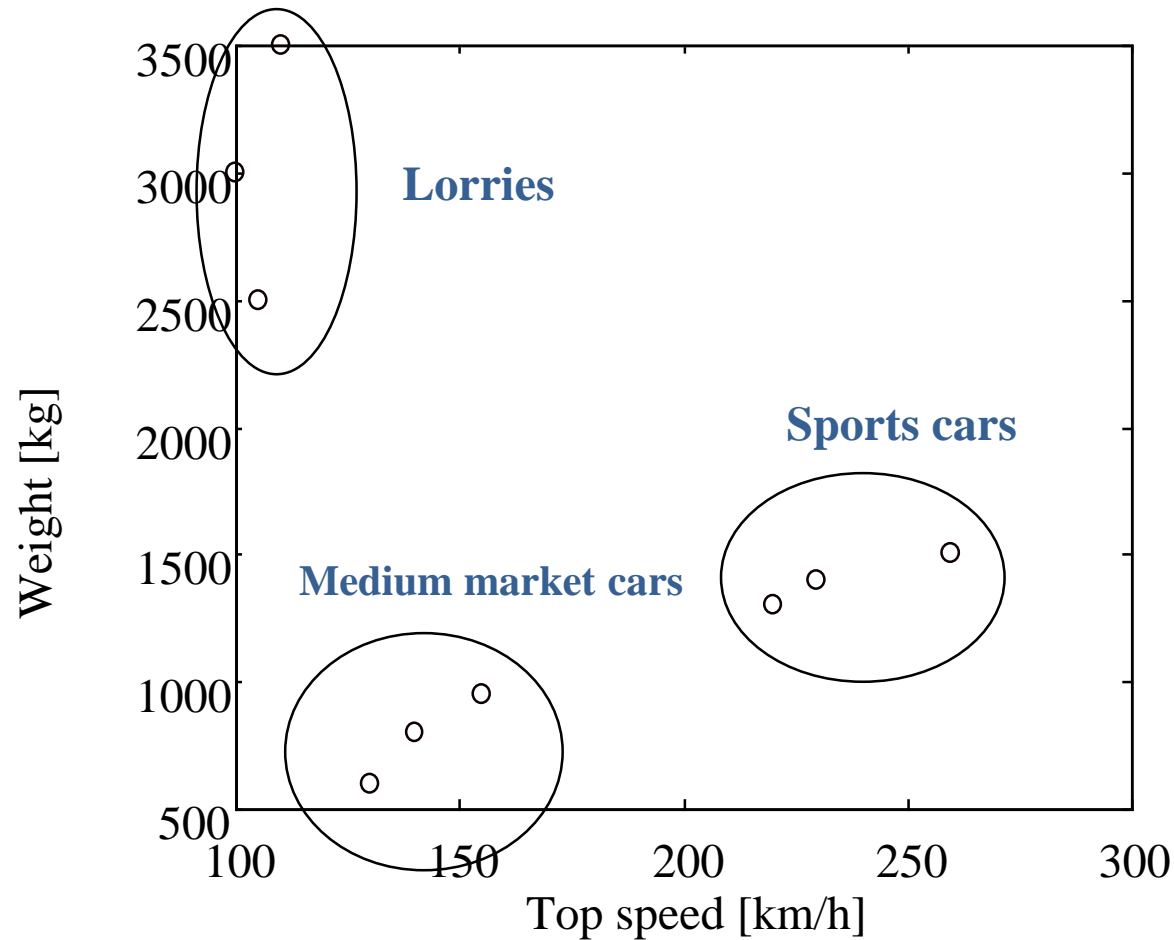
- 차량의 특성을 가지고 grouping 을 해 보자

Could see
any group ?

Vehicle	Top speed km/h	Color	Air resistance	Weight Kg
V1	220	red	0.30	1300
V2	230	black	0.32	1400
V3	260	red	0.29	1500
V4	140	gray	0.35	800
V5	155	blue	0.33	950
V6	130	white	0.40	600
V7	100	black	0.50	3000
V8	105	red	0.60	2500
V9	110	gray	0.55	3500

Summary

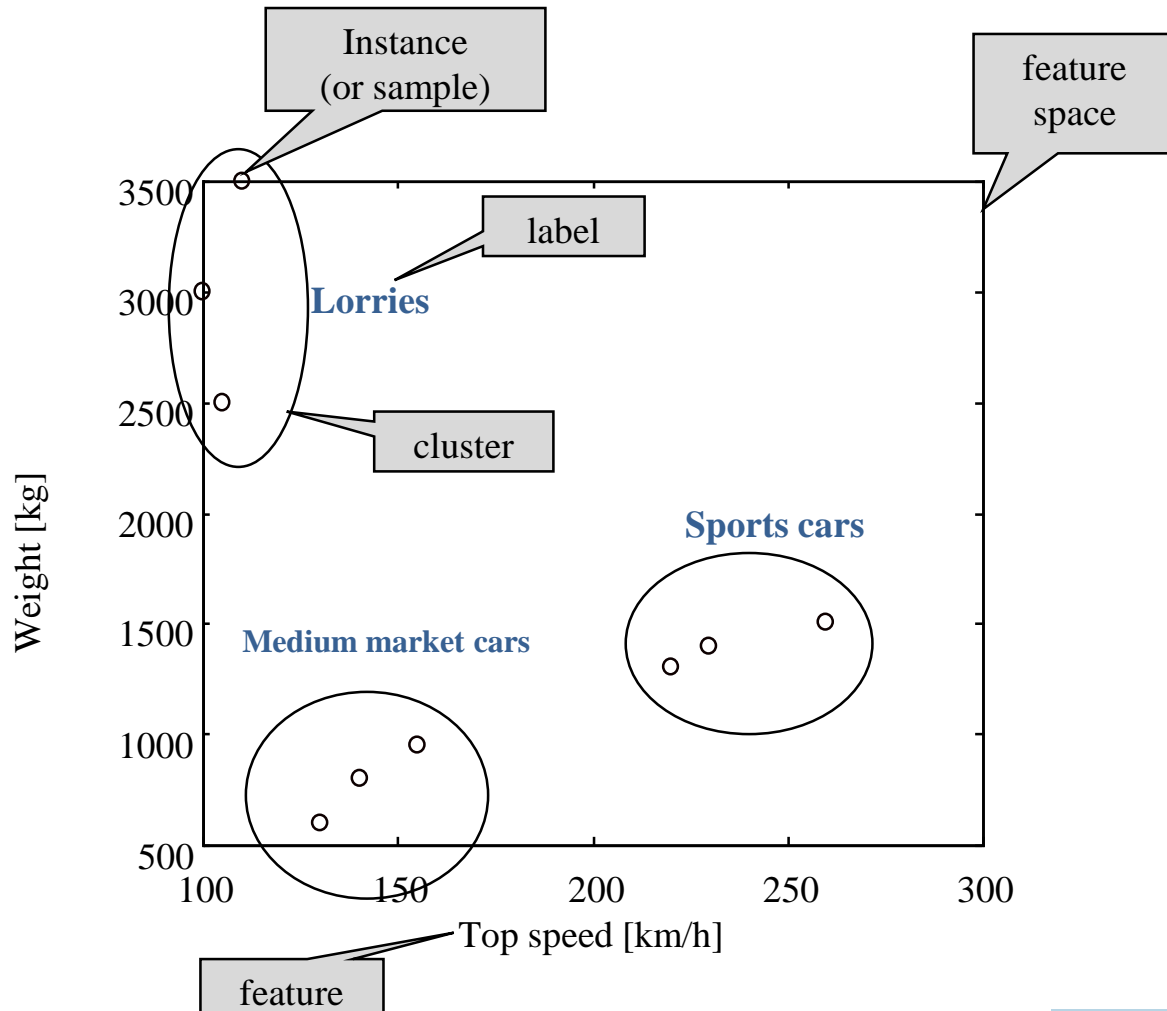
- Clustering example



Summary

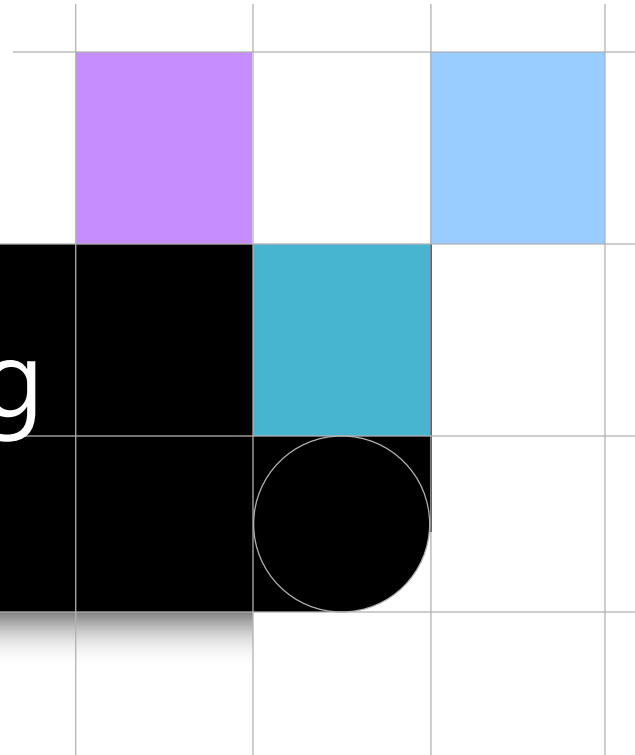
- Clustering example

Terminology



K-means clustering

(Hard c-means (HCM))



Cracked tiles example



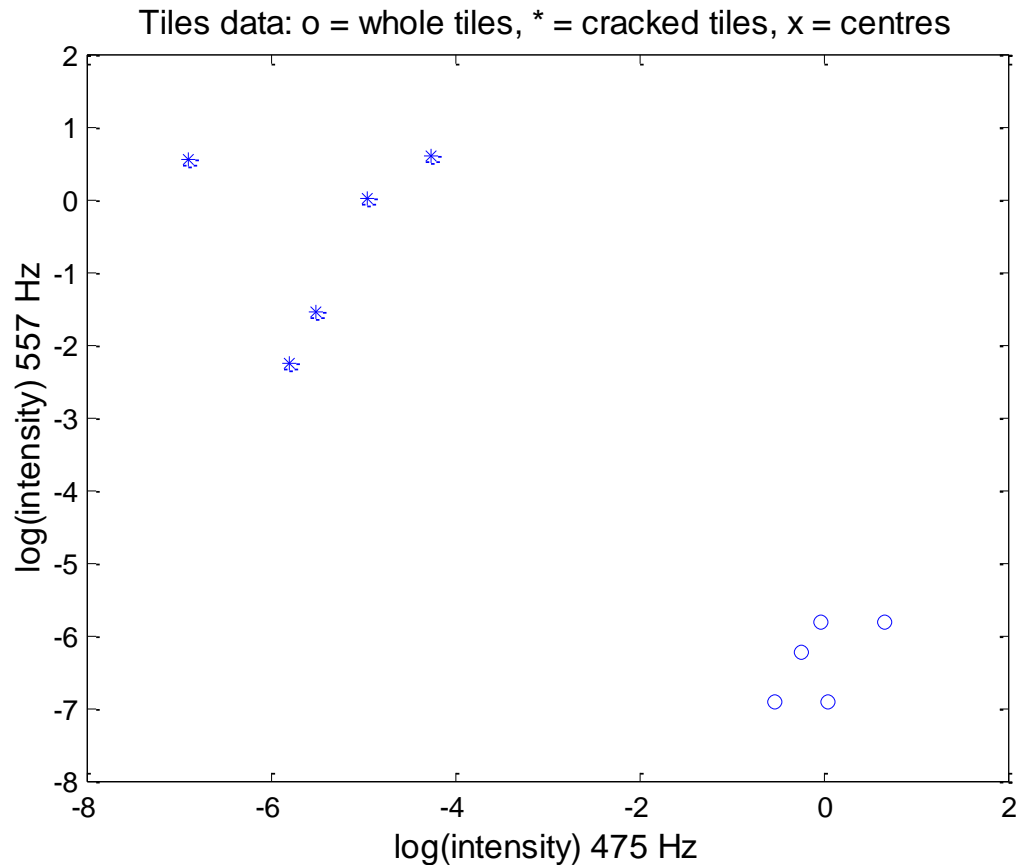
475Hz 557Hz

475Hz	557Hz
0.958	0.003
1.043	0.001
1.907	0.003
0.780	0.002
0.579	0.001
0.003	0.105
0.001	1.748
0.014	1.839
0.007	1.021
0.004	0.214

Table 1: frequency intensities for ten tiles.

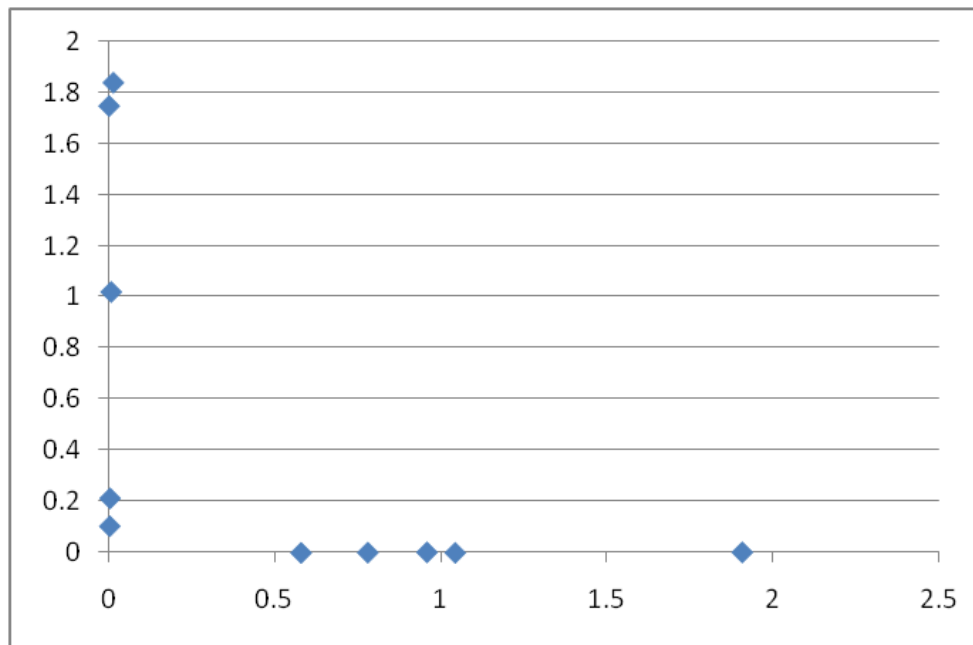
Tiles are made from clay moulded into the right shape, brushed, glazed, and baked. Unfortunately, the baking may produce invisible cracks. Operators can detect the cracks by hitting the tiles with a hammer, and in an automated system the response is recorded with a microphone, filtered, Fourier transformed, and normalised. A small set of data is given in TABLE 1 (adapted from MIT, 1997).

K-means

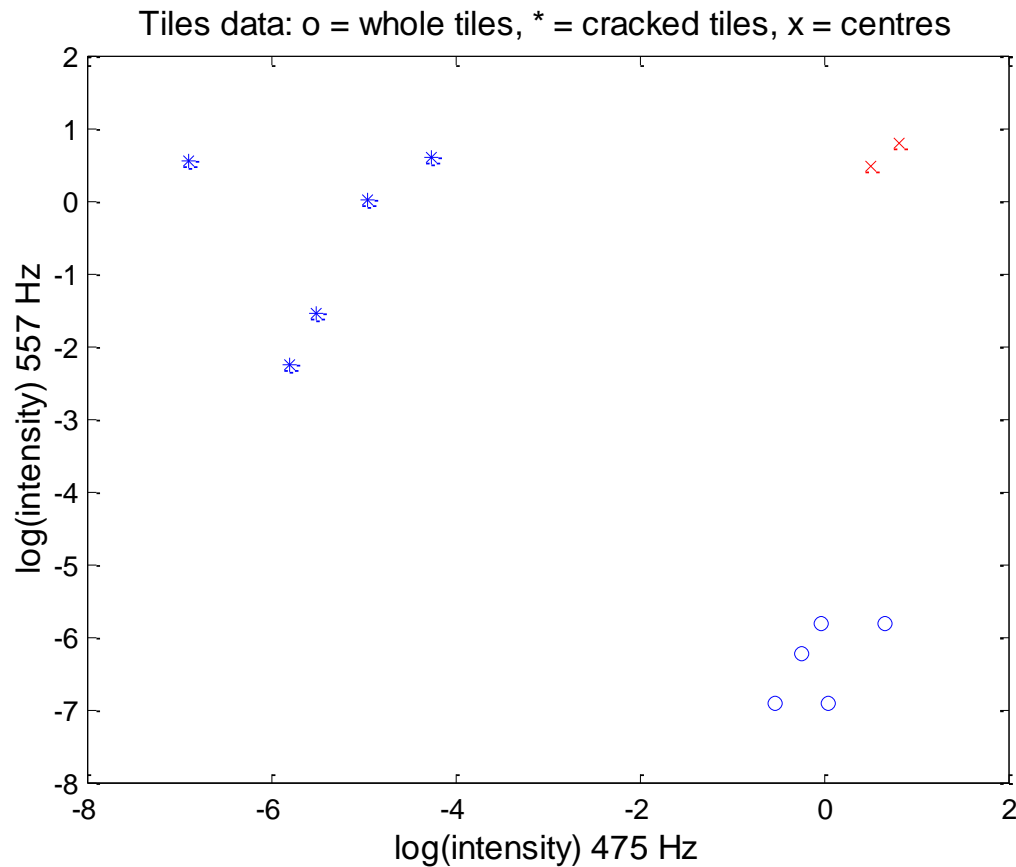


Plot of tiles by frequencies (logarithms). The whole tiles (o) seem well separated from the cracked tiles (*). The **objective** is to find the two clusters.

- Before logarithms

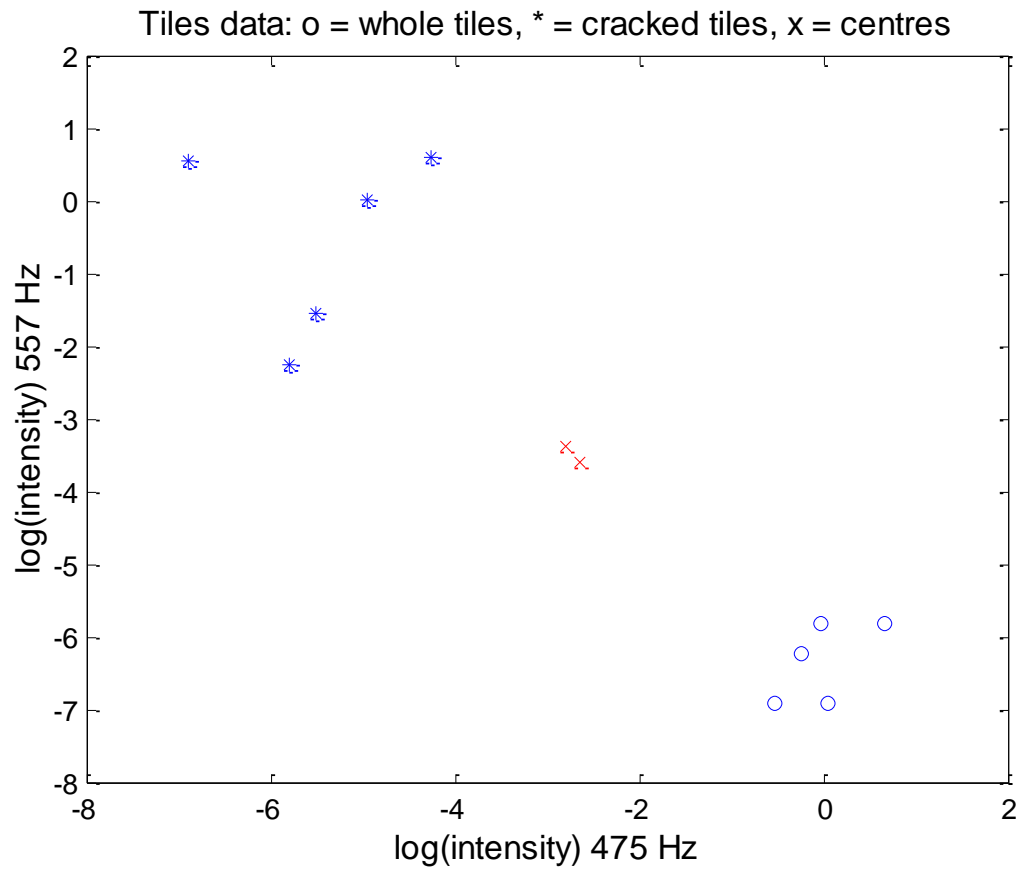


K-means



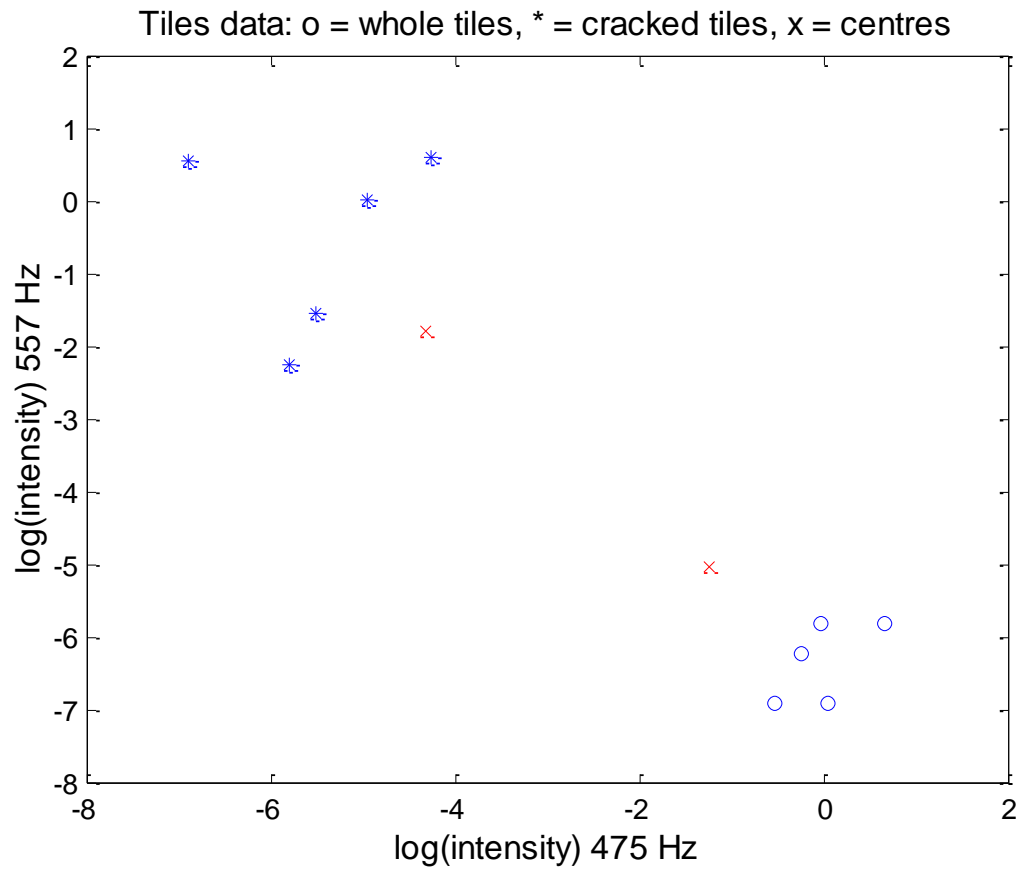
1. Place two cluster centres (x) at random.
2. Assign each data point (* and o) to the nearest cluster centre (x)

K-means



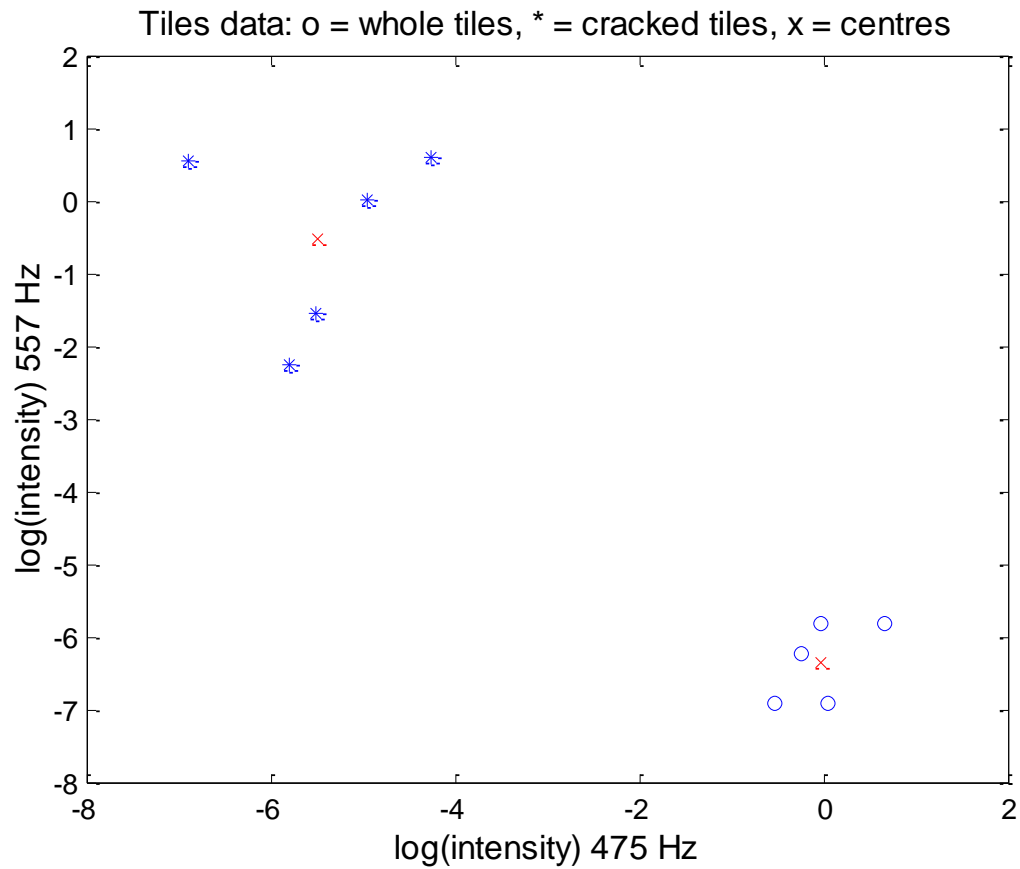
1. Compute the new centre of each class
2. Move the crosses (x)

K-means



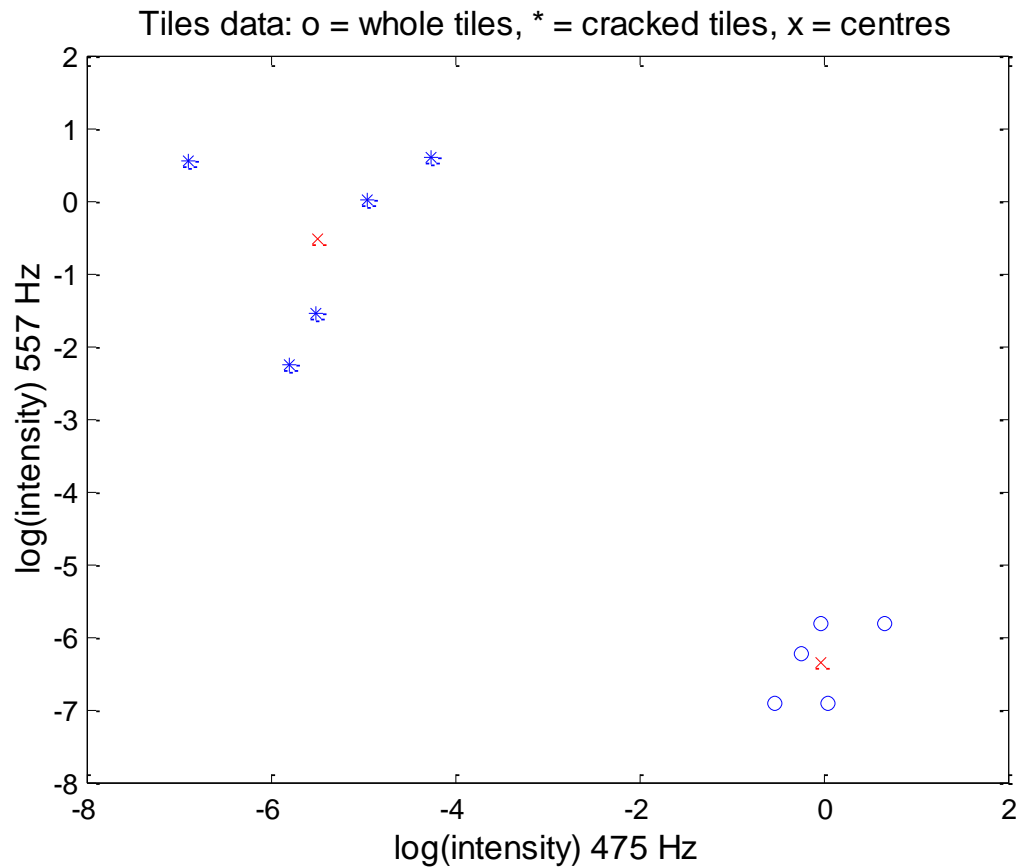
Iteration 2

K-means



Iteration 3

K-means



Iteration 4 (then stop, because no visible change)
Each data point belongs to the cluster defined by the nearest centre

K-means

475Hz 557Hz

-----+-----+

0.958 0.003

1.043 0.001

1.907 0.003

0.780 0.002

0.579 0.001

0.003 0.105

0.001 1.748

0.014 1.839

0.007 1.021

0.004 0.214

M =

0.0000 1.0000

0.0000 1.0000

0.0000 1.0000

0.0000 1.0000

0.0000 1.0000

1.0000 0.0000

1.0000 0.0000

1.0000 0.0000

1.0000 0.0000

1.0000 0.0000

First cluster

Second cluster

The membership matrix M:

1. The last five data points (rows) belong to the first cluster (column)
2. The first five data points (rows) belong to the second cluster (column)

Membership matrix **M**

data point k

cluster centre i

cluster centre j

$$m_{ik} = \begin{cases} 1 & \text{if } \|\mathbf{u}_k - \mathbf{c}_i\|^2 \leq \|\mathbf{u}_k - \mathbf{c}_j\|^2 \\ 0 & \text{otherwise} \end{cases}$$

distance

- Property of clustering

All clusters C
together fills
the whole
universe U

Clusters do not
overlap

$$\bigcup_{i=1}^c C_i = U$$

$$C_i \cap C_j = \emptyset \quad \text{for all } i \neq j$$

$$\emptyset \subset C_i \subset U \quad \text{for all } i$$

$$2 \leq c \leq K$$

A cluster C is
never empty
and it is
smaller than
the whole
universe U

There must be at
least 2 clusters in a
 c -partition and at
most as many as the
number of data
points K

Euclidean distance

$$\mathbf{p} = (p_1, p_2, p_3, \dots, p_n), \quad \mathbf{q} = (q_1, q_2, q_3, \dots, q_n)$$

Euclidean distance

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots + (p_n - q_n)^2} = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}.$$

solar

- Usage

```
kmeans(x, centers, iter.max = 10, nstart = 1,  
algorithm = c("Hartigan-Wong", "Lloyd", "Forgy",  
"MacQueen"))
```

- Argument

- **x** : numeric matrix of data, or an object that can be coerced to such a **matrix** (such as a numeric vector or a **data frame** with all numeric columns).
- **centers** : either the number of clusters, say k , or a set of initial (distinct) cluster centres.
- **iter.max** : the maximum number of iterations allowed.
- **nstart** : if centers is a number, how many random sets should be chosen? (20 또는 25 권장)
- **algorithm** : character: may be abbreviated.

- Return values
 - ◉ **cluster** : A vector of integers (from 1:k) indicating the cluster to which each point is allocated.
 - ◉ **centers** : A matrix of cluster centres.
 - ◉ **withinss** : Vector of within-cluster sum of squares, one component per cluster.
 - ◉ **tot.withinss** : Total within-cluster sum of squares, i.e. `sum(withinss)`.
 - ◉ **betweenss** : The between-cluster sum of squares, i.e. `totss-tot.withinss`.
 - ◉ **size** : The number of points in each cluster.

[R 실습] : K-means

```
require(graphics)
# a 2-dimensional example
x <- rbind(matrix(rnorm(100, sd = 0.3), ncol = 2),
            matrix(rnorm(100, mean = 1, sd = 0.3),
                  ncol = 2))
colnames(x) <- c("x", "y")
plot(x)
cl <- kmeans(x, 2)
cl # show clustering result
plot(x, col = cl$cluster)
points(cl$centers, col = 1:2, pch = 8, cex=2)

kmeans(x,1)$withinss

# random starts do help here with too many clusters
cl <- kmeans(x, 5, nstart = 25)
plot(x, col = cl$cluster)
points(cl$centers, col = 1:5, pch = 8)
```

[R 실습] : K-means

```
> cl
K-means clustering with 2 clusters of sizes 49, 51

Cluster means:
              x              y
1 -0.08673691  0.02745475
2  1.01929789  1.07596115

Clustering vector:
 [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
[37] 1 1 1 2 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
[73] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2

Within cluster sum of squares by cluster:
[1] 7.078592 8.616434
(between_SS / total_SS =  78.7 %)
```

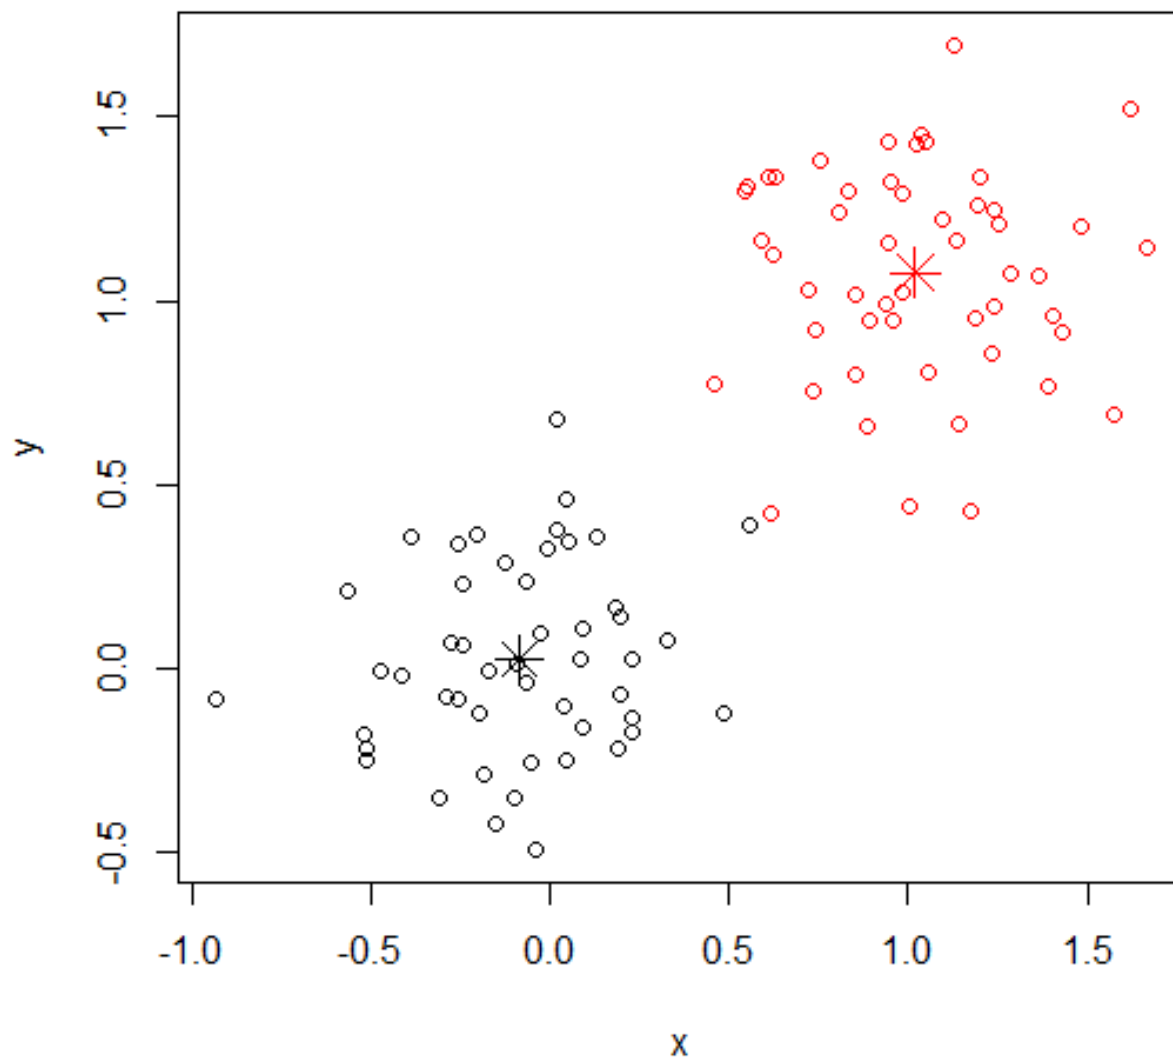
Available components:

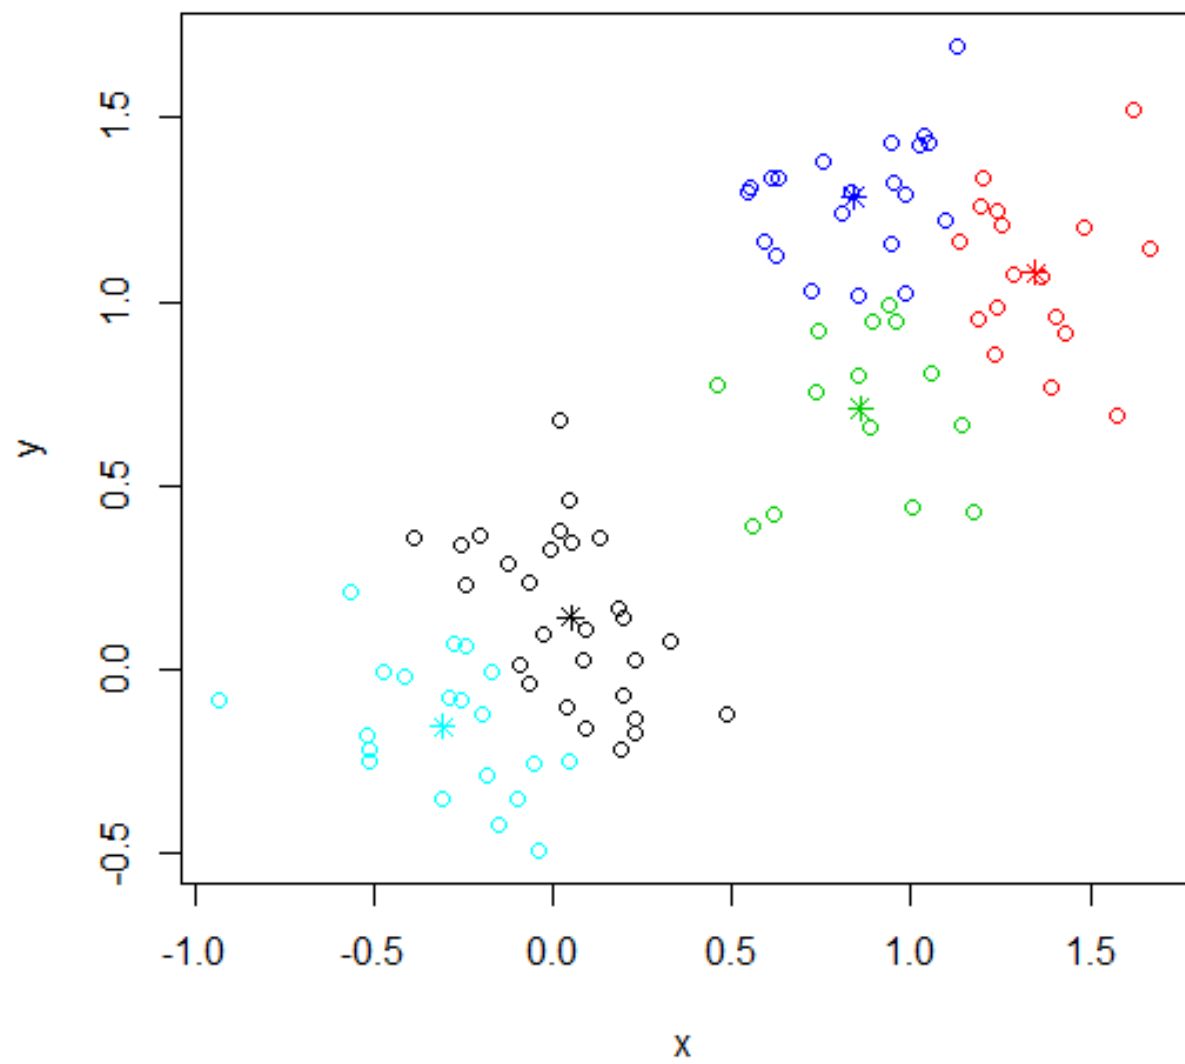
```
[1] "cluster"      "centers"      "totss"        "withinss"
[5] "tot.withinss" "betweenss"    "size"         "iter"
[9] "ifault"
```


[R 실습] : K-means

```
> cl$cluster
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
[37] 1 1 1 2 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
[73] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
> cl$centers
      x      y
1 -0.08673691 0.02745475
2  1.01929789 1.07596115
> cl$withinss
[1] 7.078592 8.616434
> cl$betweenss
[1] 58.04374
> cl$size
[1] 49 51
```

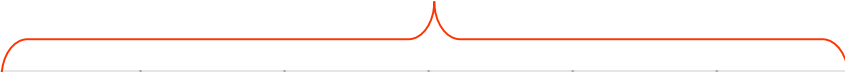
[R 실습] : K-means





[R 실습] : K-means

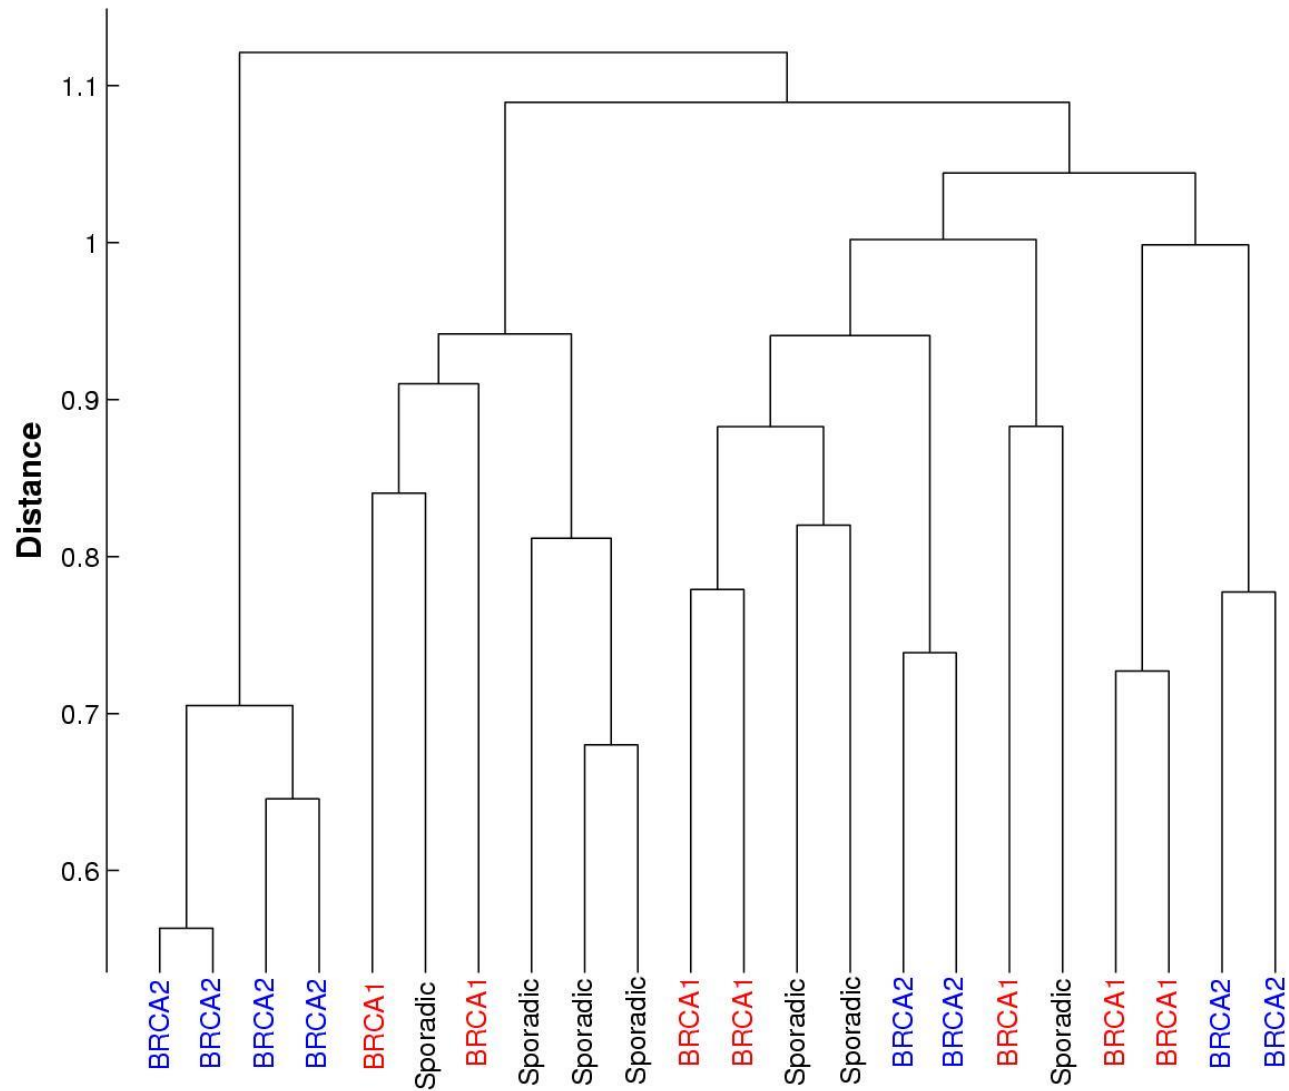
- 제공한 snsdata.csv 파일에 대해 kmeans clustering test 를 하시오
 - Remove NA rows
 - Collect data by $18 \leq \text{age} \leq 20$
 - Change gender value : M \rightarrow 1, F \rightarrow 0
 - Set k = 5



gradyear	gender	age	friends	basketball	football	soccer	softball	volleyball	swimming	cheerl
2006	M	18.982	7	0	0	0	0	0	0	0
2006	F	18.801	0	0	1	0	0	0	0	0
2006	M	18.335	69	0	1	0	0	0	0	0
2006	F	18.875	0	0	0	0	0	0	0	0
2006	NA	18.995	10	0	0	0	0	0	0	0
2006	F		142	0	0	0	0	0	0	0
2006	F	18.93	72	0	0	0	0	0	0	0
2006	M	18.322	17	0	0	0	1	0	0	0
2006	F	19.055	52	0	0	0	0	0	0	0
2006	F	18.708	39	0	0	0	0	0	0	0

- Try to find characteristics of each cluster (consider above 6 features)

Hierarchical Clustering

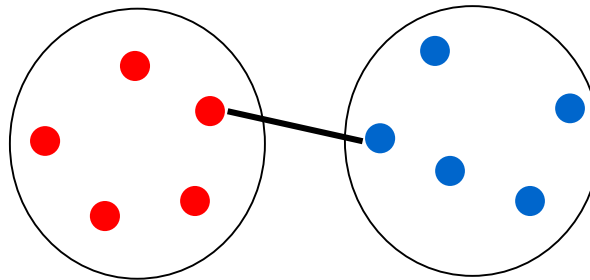


Hierarchical Clustering

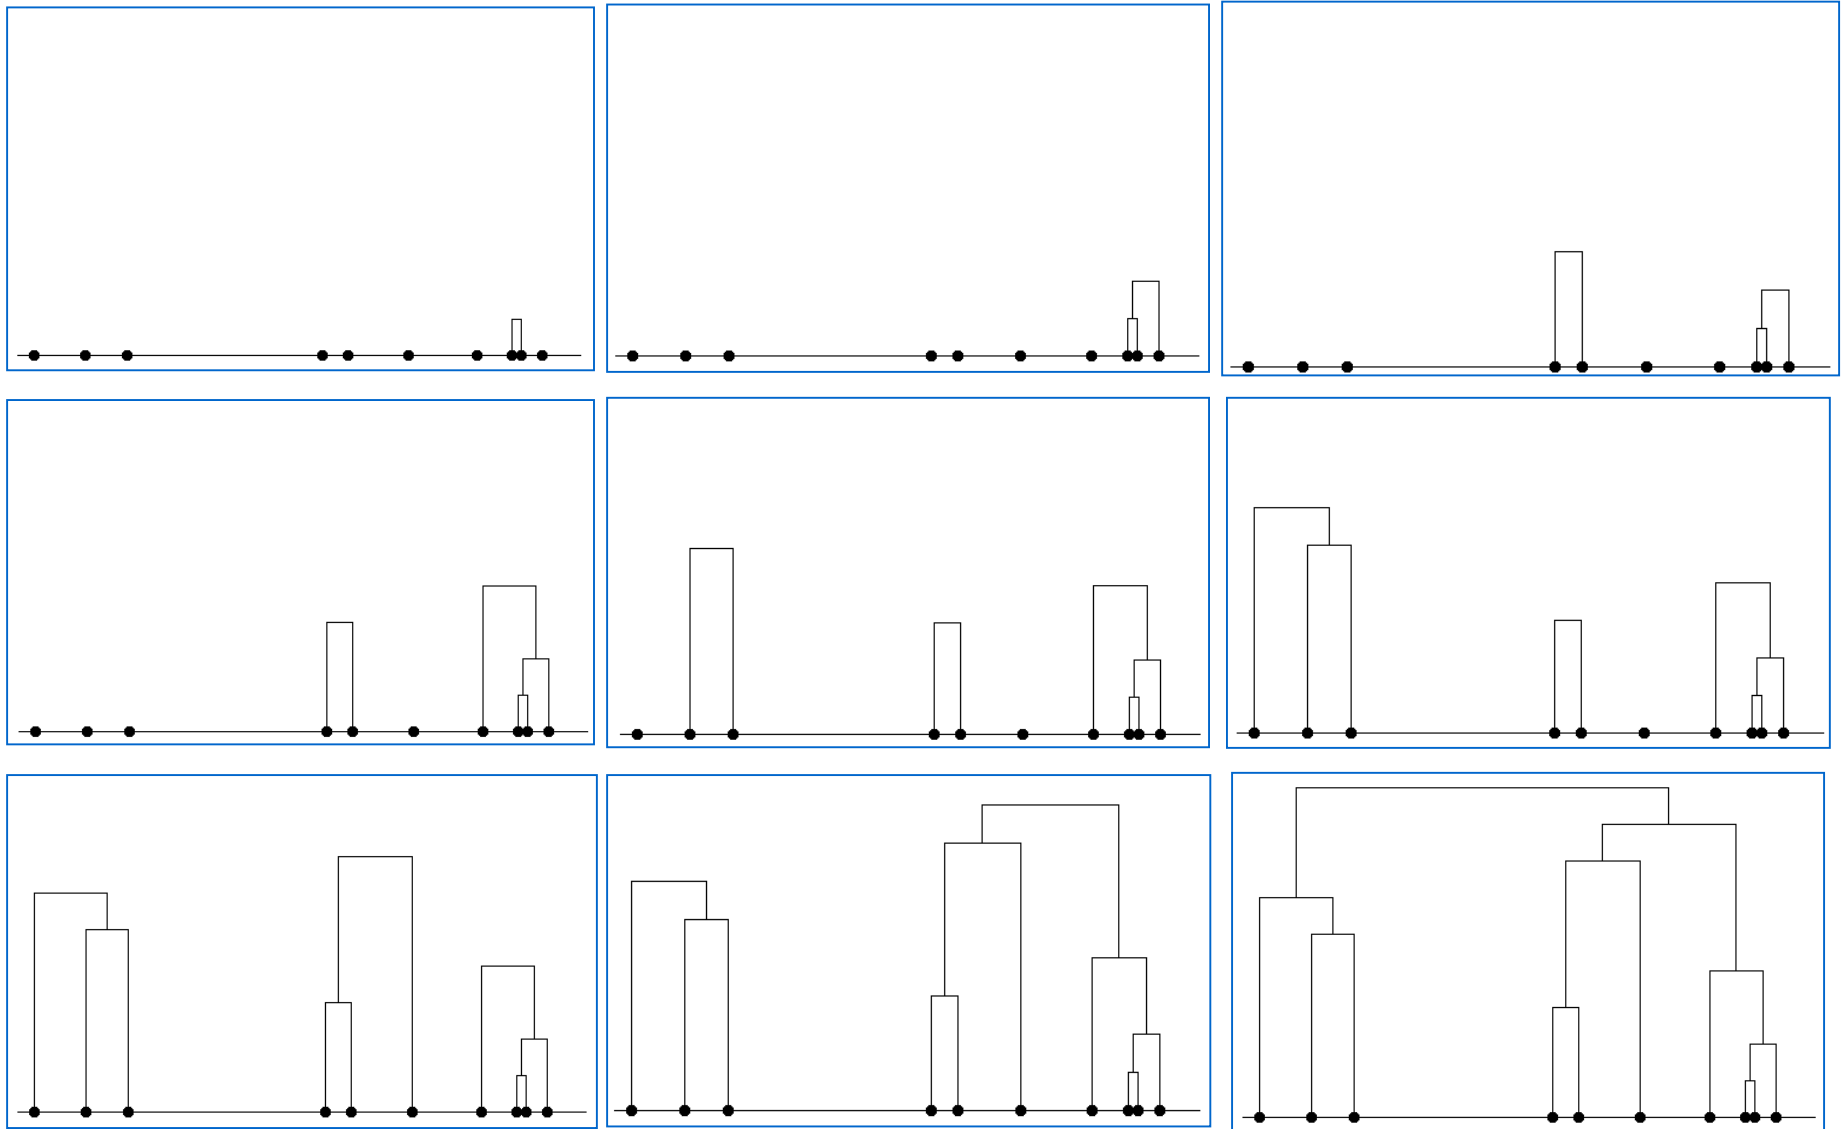
- Given a set of N items to be clustered, and an $N \times N$ distance (or similarity) matrix, the basic process of hierarchical clustering (defined by S.C. Johnson in 1967) is this:
 - ① Start by assigning each item to a cluster, so that if you have N items, you now have N clusters, each containing just one item. Let the distances (similarities) between the clusters the same as the distances (similarities) between the items they contain.
 - ② Find the closest (most similar) pair of clusters and merge them into a single cluster, so that now you have one cluster less.
 - ③ Compute distances (similarities) between the new cluster and each of the old clusters.
 - ④ Repeat steps 2 and 3 until all items are clustered into a single cluster of size N .

Hierarchical Clustering

- Step 3 can be done in different ways, which is what distinguishes **single-linkage** from **complete-linkage** and **average-linkage** clustering.
- In **single-linkage clustering** (also called the *connectedness* or *minimum* method), we consider the distance between one cluster and another cluster to be equal to the shortest distance from any member of one cluster to any member of the other cluster.
- If the data consist of similarities, we consider the similarity between one cluster and another cluster to be equal to the greatest similarity from any member of one cluster to any member of the other cluster.

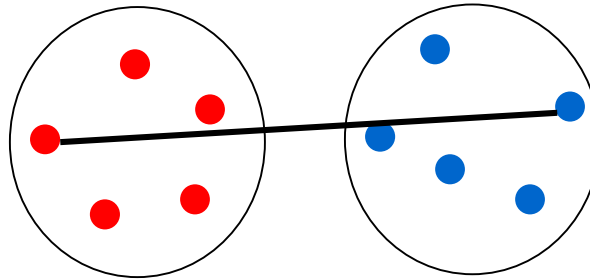


[Single-linkage]

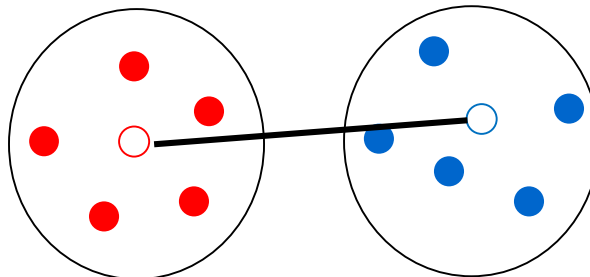


Hierarchical Clustering

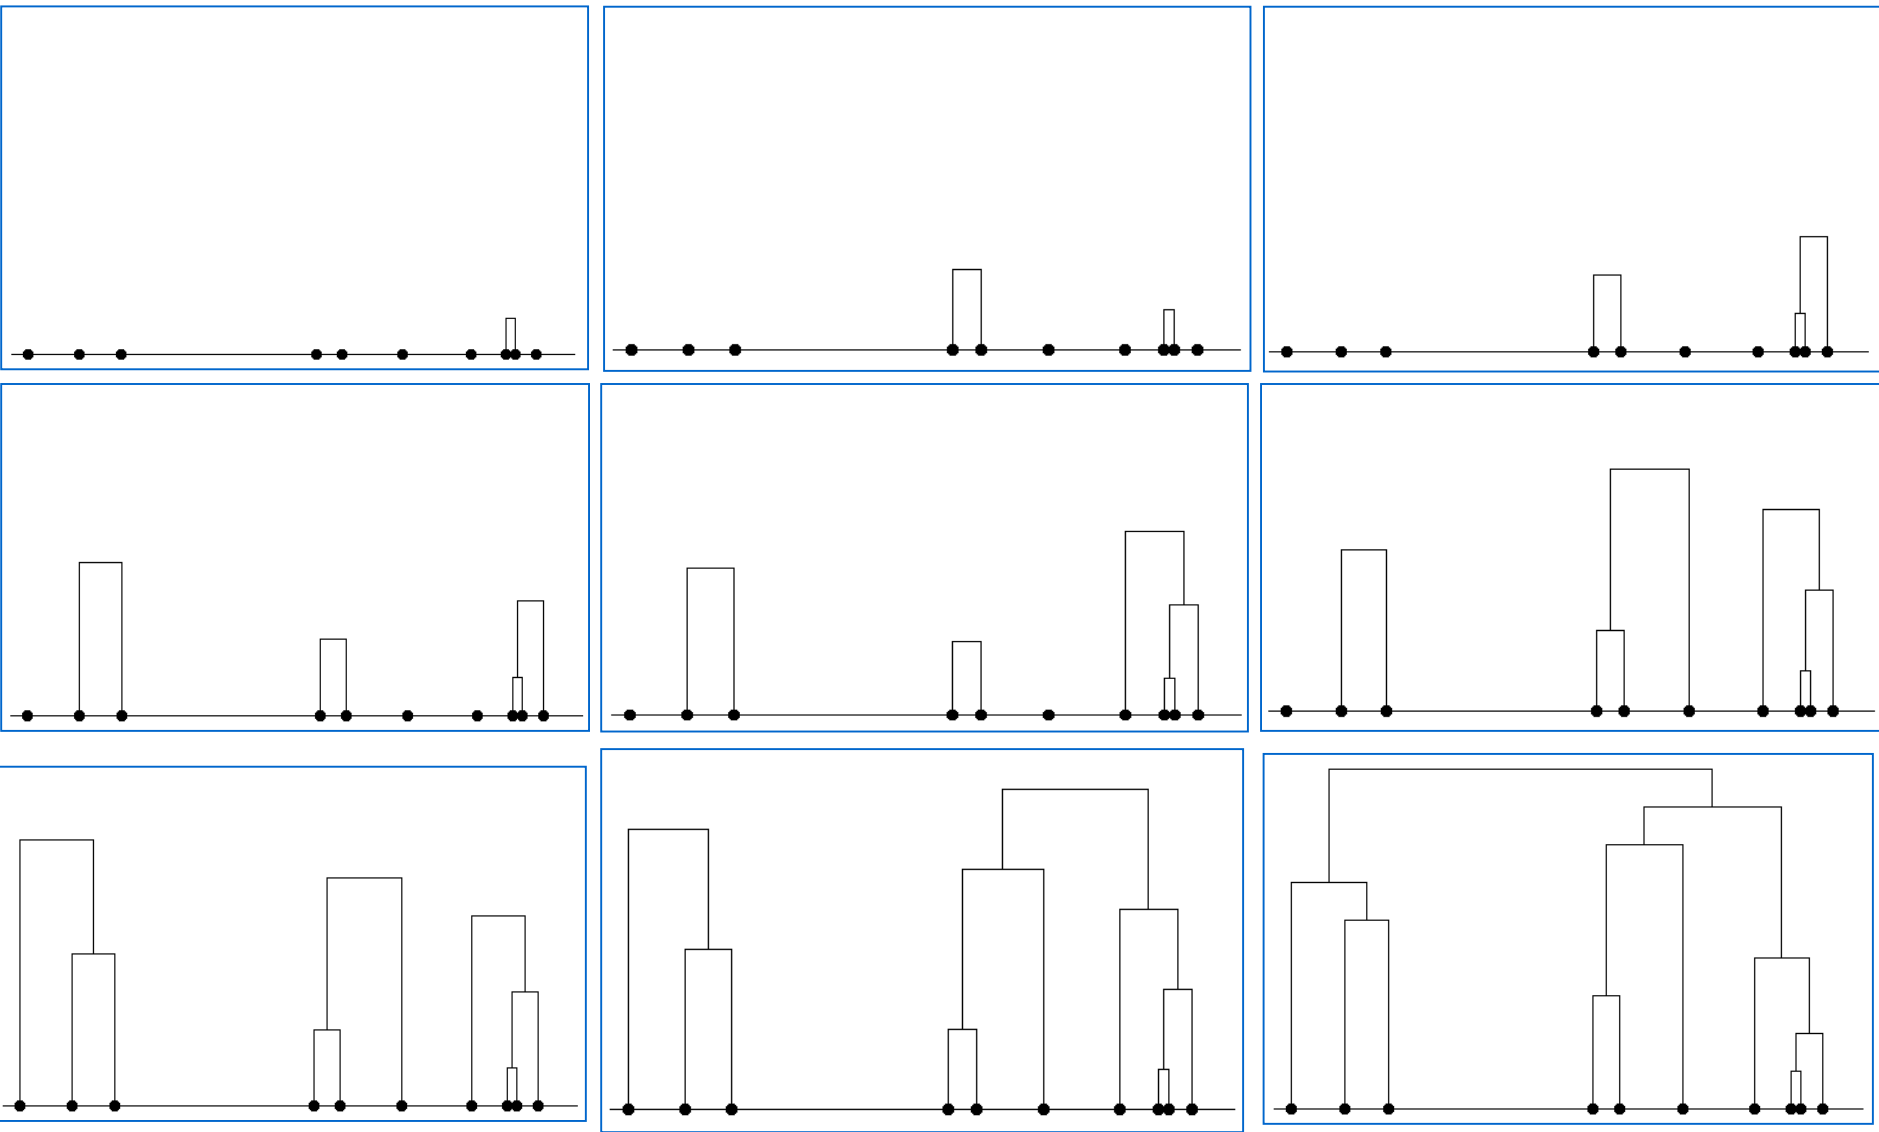
- In **complete-linkage clustering** (also called the *diameter* or *maximum* method), we consider the distance between one cluster and another cluster to be equal to the greatest distance from any member of one cluster to any member of the other cluster.



- In **average-linkage clustering**, we consider the distance between one cluster and another cluster to be equal to the average distance from any member of one cluster to any member of the other cluster.



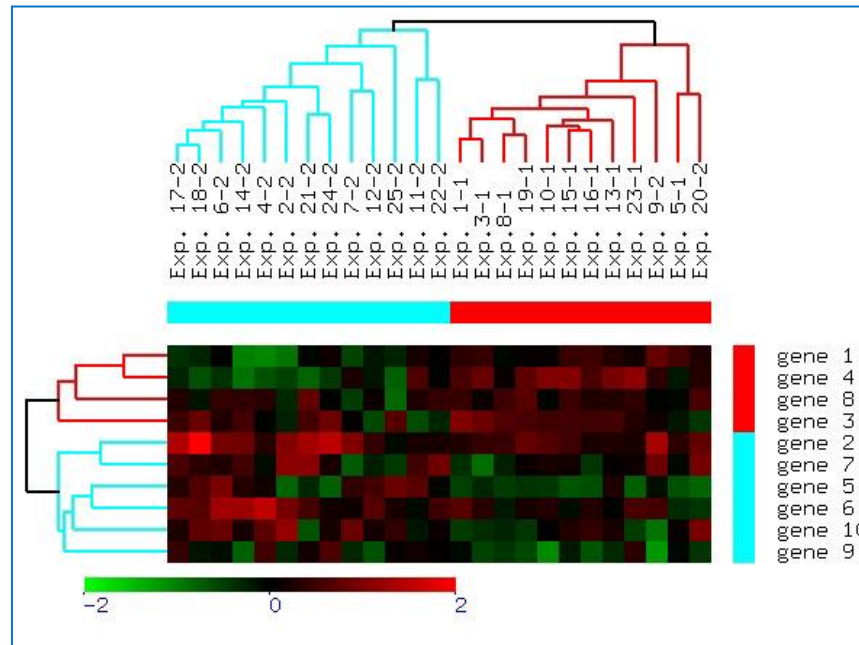
[Complete-linkage]



Hierarchical Clustering

- Problems

- The main weaknesses of agglomerative clustering methods are:
- they do not scale well: time complexity of at least $O(n^2)$, where n is the number of total objects;
- they can never undo what was done previously.

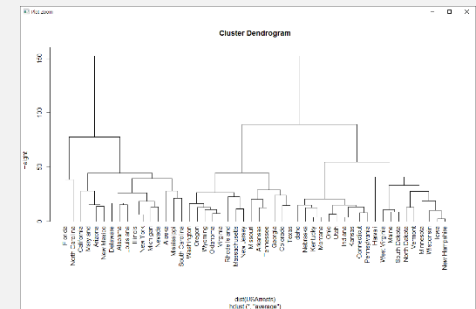


- hclust (in stats library)
- Usage

```
hclust(d, method = "complete", members=NULL)
```

- Argument
 - **d**: a dissimilarity structure as produced by dist.
 - **method** : the agglomeration method to be used. This should be (an unambiguous abbreviation of) one of "ward", "single", "complete", "average", "mcquitty", "median" or "centroid".
 - **members**: NULL or a vector with length size of d. See the ‘Details’ section

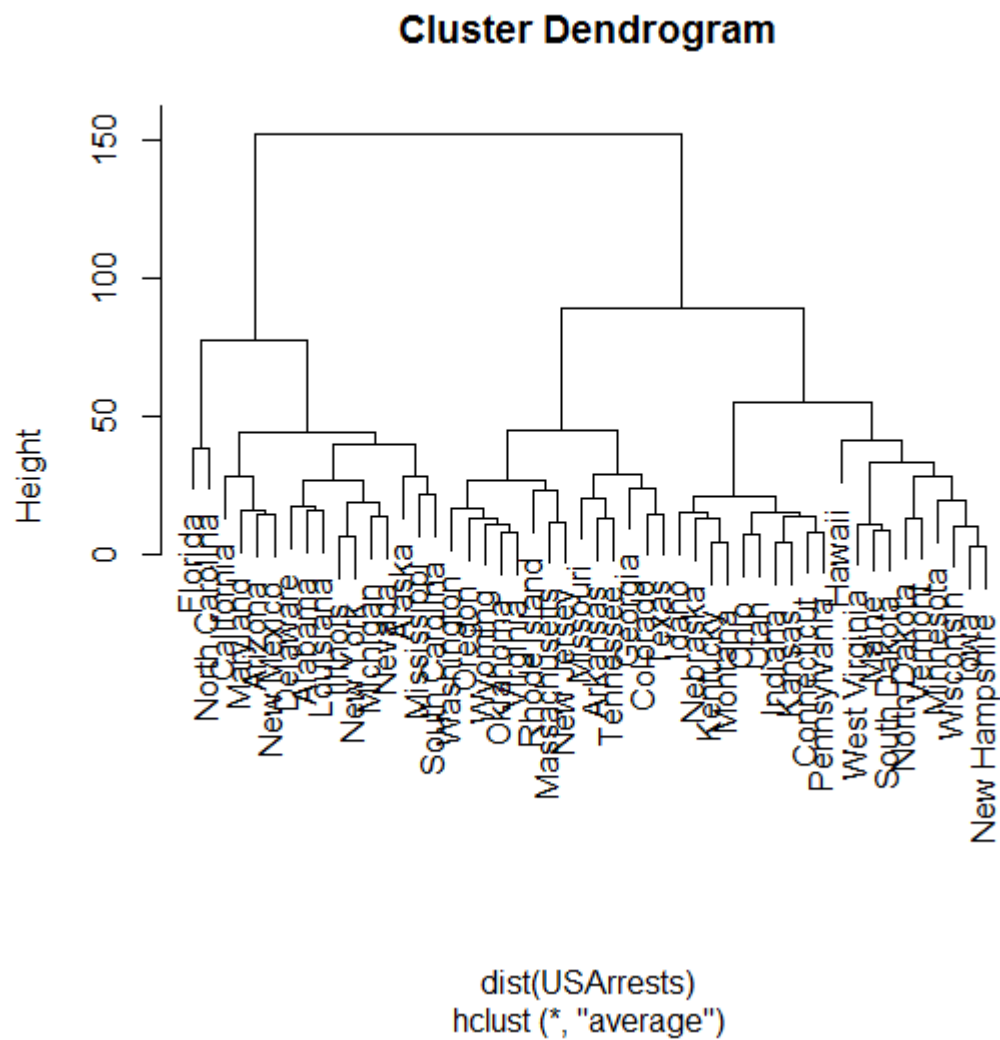
```
require(stats)
require(graphics)
hc <- hclust(dist(USArrests), "ave")
plot(hc)
plot(hc, hang = -1)
```

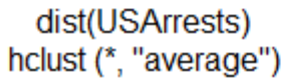


```
## Do the same with centroid clustering and squared
## Euclidean distance,
## cut the tree into ten clusters and reconstruct
## the upper part of the
## tree from the cluster centers.
```

```
hc <- hclust(dist(USArrests)^2, "cen")
memb <- cutree(hc, k = 10)
memb
```

```
> plot(hc)
```





> memb

Alabama	Alaska	Arizona	Arkansas
1	2	3	4
California	Colorado	Connecticut	Delaware
3	4	5	1
Florida	Georgia	Hawaii	Idaho
6	4	7	5
Illinois	Indiana	Iowa	Kansas
1	5	8	5
Kentucky	Louisiana	Maine	Maryland
5	1	8	3
Massachusetts	Michigan	Minnesota	Mississippi
9	1	8	2
Missouri	Montana	Nebraska	Nevada
4	5	5	1
New Hampshire	New Jersey	New Mexico	New York
8	9	3	1
North Carolina	North Dakota	Ohio	Oklahoma
10	8	5	9
Oregon	Pennsylvania	Rhode Island	South Carolina
9	5	9	2
South Dakota	Tennessee	Texas	Utah
8	4	4	5
Vermont	Virginia	Washington	West Virginia
8	9	9	8
Wisconsin	Wyoming		
8	9		

10개 클러스터의 중심점 계산

```
cent <- NULL
```

```
for(k in 1:10){
```

```
  cent <- rbind(cent, colMeans(USArrests[memb == k, ,  
                                drop = FALSE]))
```

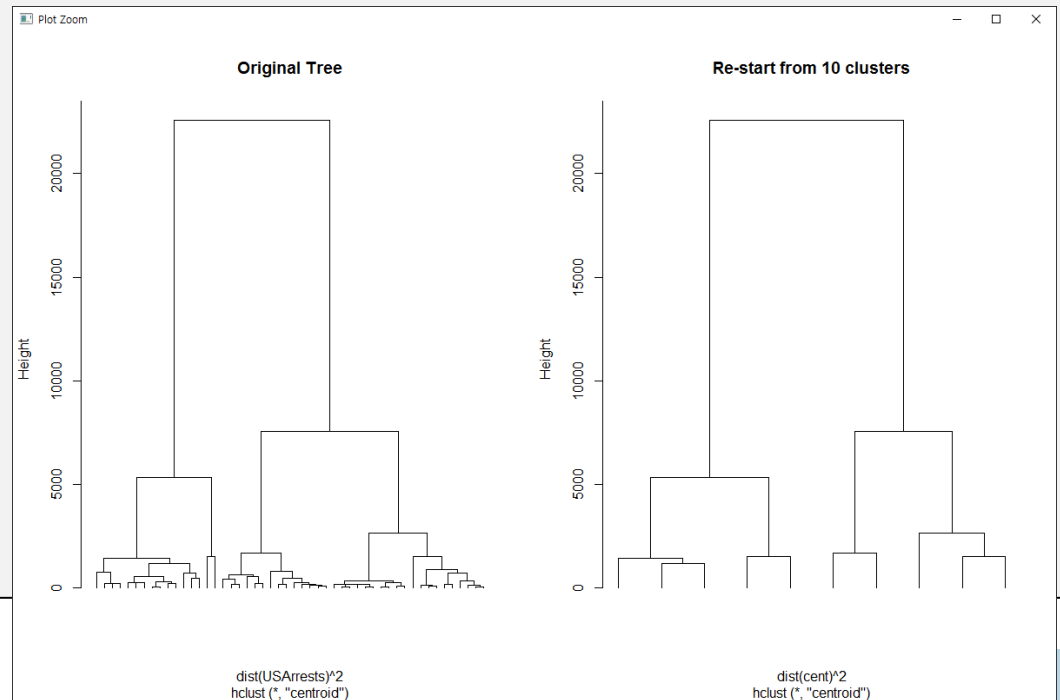
```
}
```

```
> cent
```

	Murder	Assault	UrbanPop	Rape
[1,]	11.471429	247.57143	74.28571	27.20000
[2,]	13.500000	267.00000	46.66667	28.03333
[3,]	9.950000	288.75000	77.00000	32.87500
[4,]	11.500000	195.33333	66.16667	27.43333
[5,]	5.590000	112.40000	65.60000	17.27000
[6,]	15.400000	335.00000	80.00000	31.90000
[7,]	5.300000	46.00000	83.00000	20.20000
[8,]	2.688889	64.55556	50.66667	10.54444
[9,]	5.750000	156.75000	74.00000	19.40000
[10,]	13.000000	337.00000	45.00000	16.10000

10개 클러스터를 재 클러스터링

```
hc1 <- hclust(dist(cent)^2, method = "cen",  
             members = table(memb))  
opar <- par(mfrow = c(1, 2))  
plot(hc, labels = FALSE, hang = -1,  
     main = "Original Tree")  
plot(hc1, labels = FALSE, hang = -1,  
     main = "Re-start from 10 clusters")  
par(opar)
```



[실습문제]

- UCI machine learning repository 에서 wine dataset 을 다운받아 hierarchical clustering 을 테스트 하시오.
 - <http://archive.ics.uci.edu/ml/>
 - 첫번째 컬럼은 class data 이니 클러스터링에서 제외
- 각 컬럼의 데이터 범위가 각기 달라서 거리 계산시 문제가 있음
 - 따라서 각 컬럼의 데이터 범위가 0~1 사이가 되도록 조정
- 1) **plot(hc, hang = -1)** 를 이용하여 그래프 작성
- 2) Tree 를 3개의 cluster 가 되도록 잘라서 그래프를 그린다
- 3) 3개로 클러스터링 된 결과를 실제 class 와 비교하여 일치도가 얼마나 되는지를 보이시오